

Windows Internals

Module 4: System Architecture (Part 2)

Pavel Yosifovich

CTO, CodeValue

pavel@codevalue.net

<http://blogs.Microsoft.co.il/blogs/pavely>



Contents

- Core system files
- Multiprocessing
- Subsystems and NTDLL
- System processes
- Wow64
- Summary

Core system files

- **Ntoskrnl.exe**
 - Executive and kernel on 64 bit systems
- **NtKrnlPa.exe**
 - Executive and kernel on 32 bit systems
- **Hal.dll**
 - Hardware Abstraction Layer
- **Win32k.sys**
 - Kernel component of the Windows subsystem
 - Handles windowing and GDI
- **NtDll.dll**
 - System support routines and Native API dispatcher to executive services
- **Kernel32.dll, user32.dll, gdi32.dll, advapi32.dll**
 - Core Windows subsystem DLLs
- **CSRSS.exe (“Client Server Runtime SubSystem”)**
 - The Windows subsystem process

Demo

Core system files

Symmetric multiprocessing

- **SMP**
 - All CPUs are the same and share main memory and have equal access to peripheral devices (no master/slave)
- **Basic architecture supports up to 32/64 CPUs**
 - Windows 7 64 bit & 2008 R2 support up to 256 cores
 - Uses a new concept of a “processor group”
- **Actual number of CPUs determined by licensing and product type**
 - Multiple cores do not count towards this limit

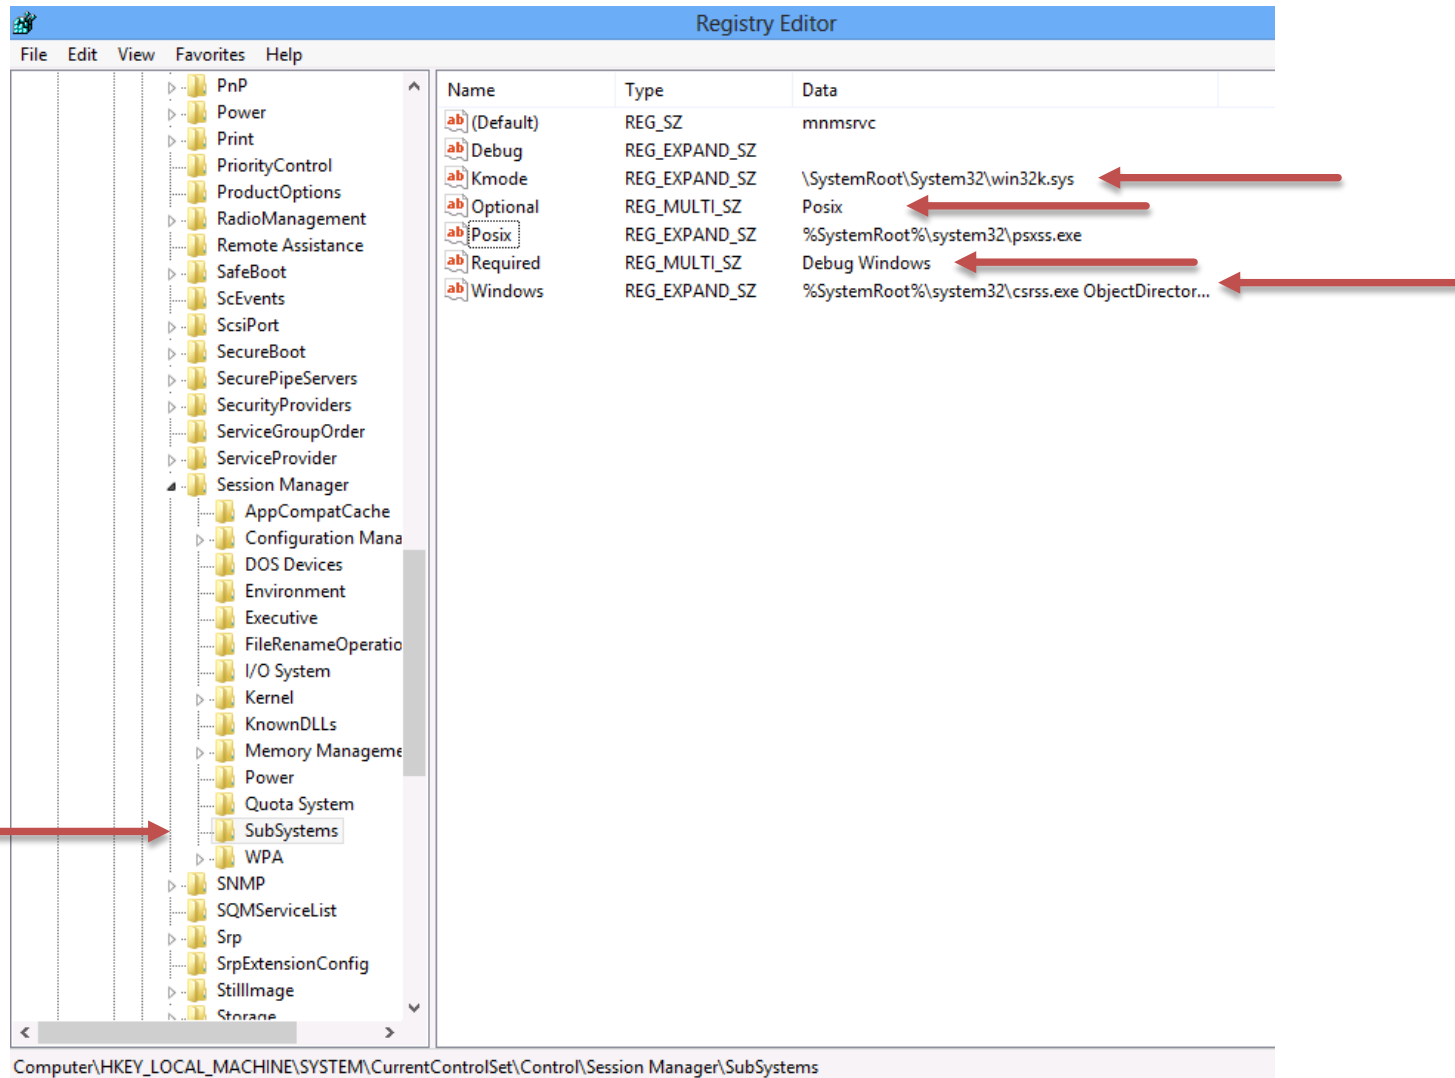
Demo

SMP

Subsystems

- **A subsystem is a special view of the OS**
 - Exposes services via subsystem DLLs
- **Original NT shipped with Win32, OS/2 and POSIX 1003.1 (POSIX-1)**
 - Windows XP dropped support for OS/2
 - An enhanced POSIX version is available with the “Services for UNIX” product
- **The Windows subsystem must always be running**
 - Owner of keyboard, mouse and display
- **Some API functions use the Advanced Local Procedure Call (ALPC) to notify CSRSS of relevant events**
- **Other subsystems configured to load on demand**
- **Subsystem information stored in registry:**
HKLM\System\CCS\Control\Session Manager\Subsystems

Subsystems in the registry



Subsystem DLLs

- **Every image belongs to exactly one subsystem**
 - Value stored in image PE header
 - Can view with Dependency Walker (**depends.exe**)
 - Allows the Windows Loader to make correct decisions
- **An image of a certain subsystem calls API functions exposed through the subsystem DLLs**
 - E.g. kernel32.dll, user32.dll, etc. for the Windows subsystem
- **Some images belong to no subsystem**
 - “Native” images
 - Which API functions do they call?

The Native API

- **Implemented by NTDLL.DLL**
 - Used by subsystem DLLs and “native” images
 - Undocumented interface
 - Lowest layer of user mode code
- **Contains**
 - Various support functions
 - Dispatcher to kernel services
 - Most of them accessible using Windows API “wrappers”

Demo

Subsystem DLLs and NTDLL

System Processes

- **Idle process**
- **System process**
- **Session Manager (Smss.Exe)**
- **Windows subsystem (Csrss.Exe)**
- **Logon process (Winlogon.Exe)**
- **Service control manager (SCM) (Services.Exe)**
- **Local security authentication server (Lsass.Exe)**
- **Local session manager (Lsm.exe)**

Idle Process

- Always has a PID of 0
- Not a real process (does not run any executable image)
- One thread per CPU (core)
- Accounts for idle time

System Process

- Has a fixed PID (4)
- Represents the kernel address space and resources
- Hosts system threads
 - Threads created by the kernel and device drivers
 - Execute code in system space only
 - Created using the **PsCreateSystemThread** kernel API (documented in the WDK)
 - Allocate memory from the system pools

Demo

Idle and System processes

Session Manager

- Running the image `\windows\system32\smss.exe`
- The first user mode process created by the system
- **Main tasks**
 - Creating system environment variables
 - Launches the subsystem processes (normally just csrss.exe)
 - Launches itself in other sessions
 - That instance loads WINLOGON and CSRSS in that session
 - Then terminates
- **Finally**
 - Waits forever for csrss.exe instances to terminate
 - If any of them dies, crashes the system
 - Waits for subsystem creation requests
 - Waits for terminal services session creation requests

Winlogon

- Running the image `\windows\system32\winlogon.exe`
- Handles interactive logons and logoffs
- If terminated, logs off the user session
- Notified of a user request by the Secure Attention Sequence (SAS), typically Ctrl+Alt+Del
- Authenticates the user by presenting a username / password dialog (through LogonUI.exe)
 - Can be replaced
- Sends captured username and password to LSASS
 - If successfully authenticated, initiates the user's session

LSASS

- Running the image `\windows\system32\lsass.exe`
- Calls the appropriate authentication package
- Upon successful authentication, creates a token representing the user's security profile
- Returns information to Winlogon

Service Control Manager (SCM)

- Running the image `\windows\system32\services.exe`
- Responsible for starting, stopping and interacting with service processes
- **Services**
 - Similar to UNIX “daemon processes”
 - Normal Windows executables, that interact with the SCM
 - Can be started automatically when the system starts up without an interactive logon
 - Can run under “special” accounts
 - LocalSystem, NetworkService, LocalService

Local Session Manager

- Introduced in Windows Vista
 - Running the image **\windows\system32\lsm.exe**
- In windows 8, turned into a service
 - Implemented in **\windows\system32\lsm.dll**
 - Hosted in a standard svchost.exe
- Manages terminal sessions on the local machine
- Communicates requests to SMSS

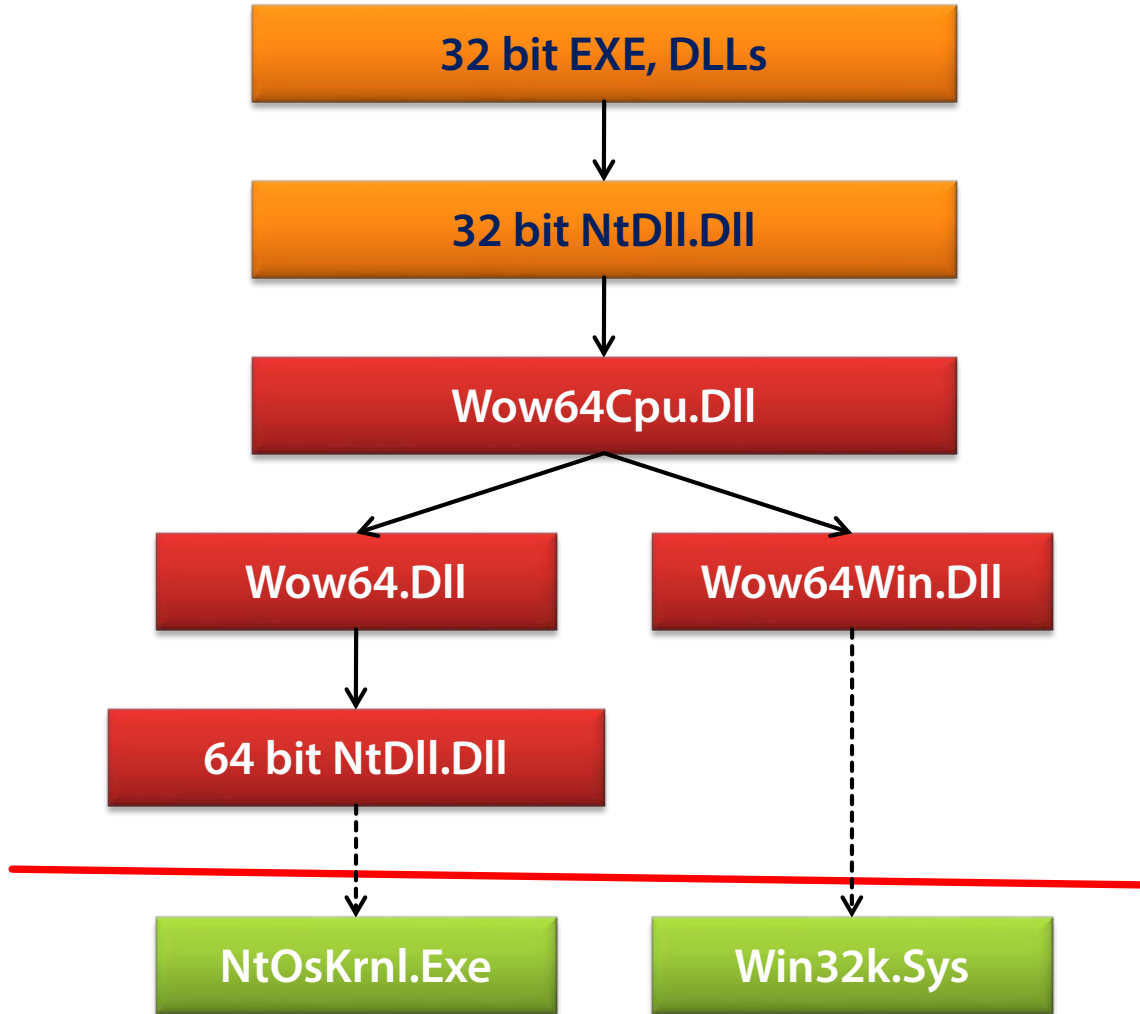
Demo

Services

Wow64

- **Allows execution of Win32 binaries on 64-bit Windows**
 - Wow64 intercepts system calls from the 32-bit application
 - Converts 32-bit data structure into 64-bit aligned structures
 - Issues the native 64-bit system call
 - Returns any data from the 64-bit system call
- **The `IsWow64Process` function can tell whether a process is running under Wow64**
- **Address space is 2GB or 4GB (if image is linked with the LARGEADDRESSAWARE flag)**
- **Device drivers must be native 64 bit**
- **File system**
 - `\windows\system32` contains 64 bit images
 - `\windows\syswow64` contains 32 bit images

Wow64 Architecture



Wow64 Restrictions

- **A 64 bit process cannot load a 32 bit DLL and vice versa**
 - Except resource-only DLLs, which can be loaded cross-architecture
- **Some APIs are not supported by Wow64 processes**
 - **E.g. ReadFileScatter, WriteFileGather**, AWE functions

File System Redirection

- System directories names have not changed in 64 bit Windows (e.g. \Windows\System32 contains native 64 bit images)
- 32 bit applications must use their own directories
 - \Windows\System32 maps to \Windows\Syswow64
 - 32 bit apps installed in \Program Files (x86)
 - 64 bit apps installed in \Program Files
- Some directories are not redirected

Registry Redirection

- Components trying to register as 32 bit and 64 bit will clash
- 32 bit components are redirected to the Wow64 registry node (**Wow6432Node**)
 - **HKEY_LOCAL_MACHINE\Software**
 - **HKEY_CLASSES_ROOT**
 - **HKEY_CURRENT_USER\Software\Classes**
- New flags for Registry APIs allow access to the 64 bit or 32 bit nodes
 - **KEY_WOW64_64KEY** – open a 64 bit key
 - **KEY_WOW64_32KEY** – open a 32 bit key

Demo

Wow64

Summary

- A Process executes under a specific subsystem
- The primary subsystem is the Windows subsystem
- NTDLL is the gateway to kernel mode
- Wow64 allows running 32 bit processes on 64 bit systems transparently