

# Windows Internals

## Module 6: Processes & Threads (Part 2)

Pavel Yosifovich

CTO, CodeValue

[pavel@codevalue.net](mailto:pavel@codevalue.net)

<http://blogs.Microsoft.co.il/blogs/pavel>



# Thread Stacks

- **Every user mode thread has two stacks**

- In kernel space (12K (x86), 24K (x64))
  - Resides in physical memory (most of the time)
- In user space (may be large)
  - By default 1MB is reserved, 64KB committed
  - A guard page is placed just below the last committed page, so that the stack can grow
  - Can change the initial size
    - Using linker settings as new defaults
    - On a thread by thread basis in the call to **CreateThread** / **CreateRemoteThread(Ex)**
    - Can specify a new committed or reserved size, but not both
    - Committed is assumed, unless the flag **STACK\_SIZE\_PARAM\_IS\_A\_RESERVATION** is used


**Demo**

**Thread stacks**

# Thread Priorities

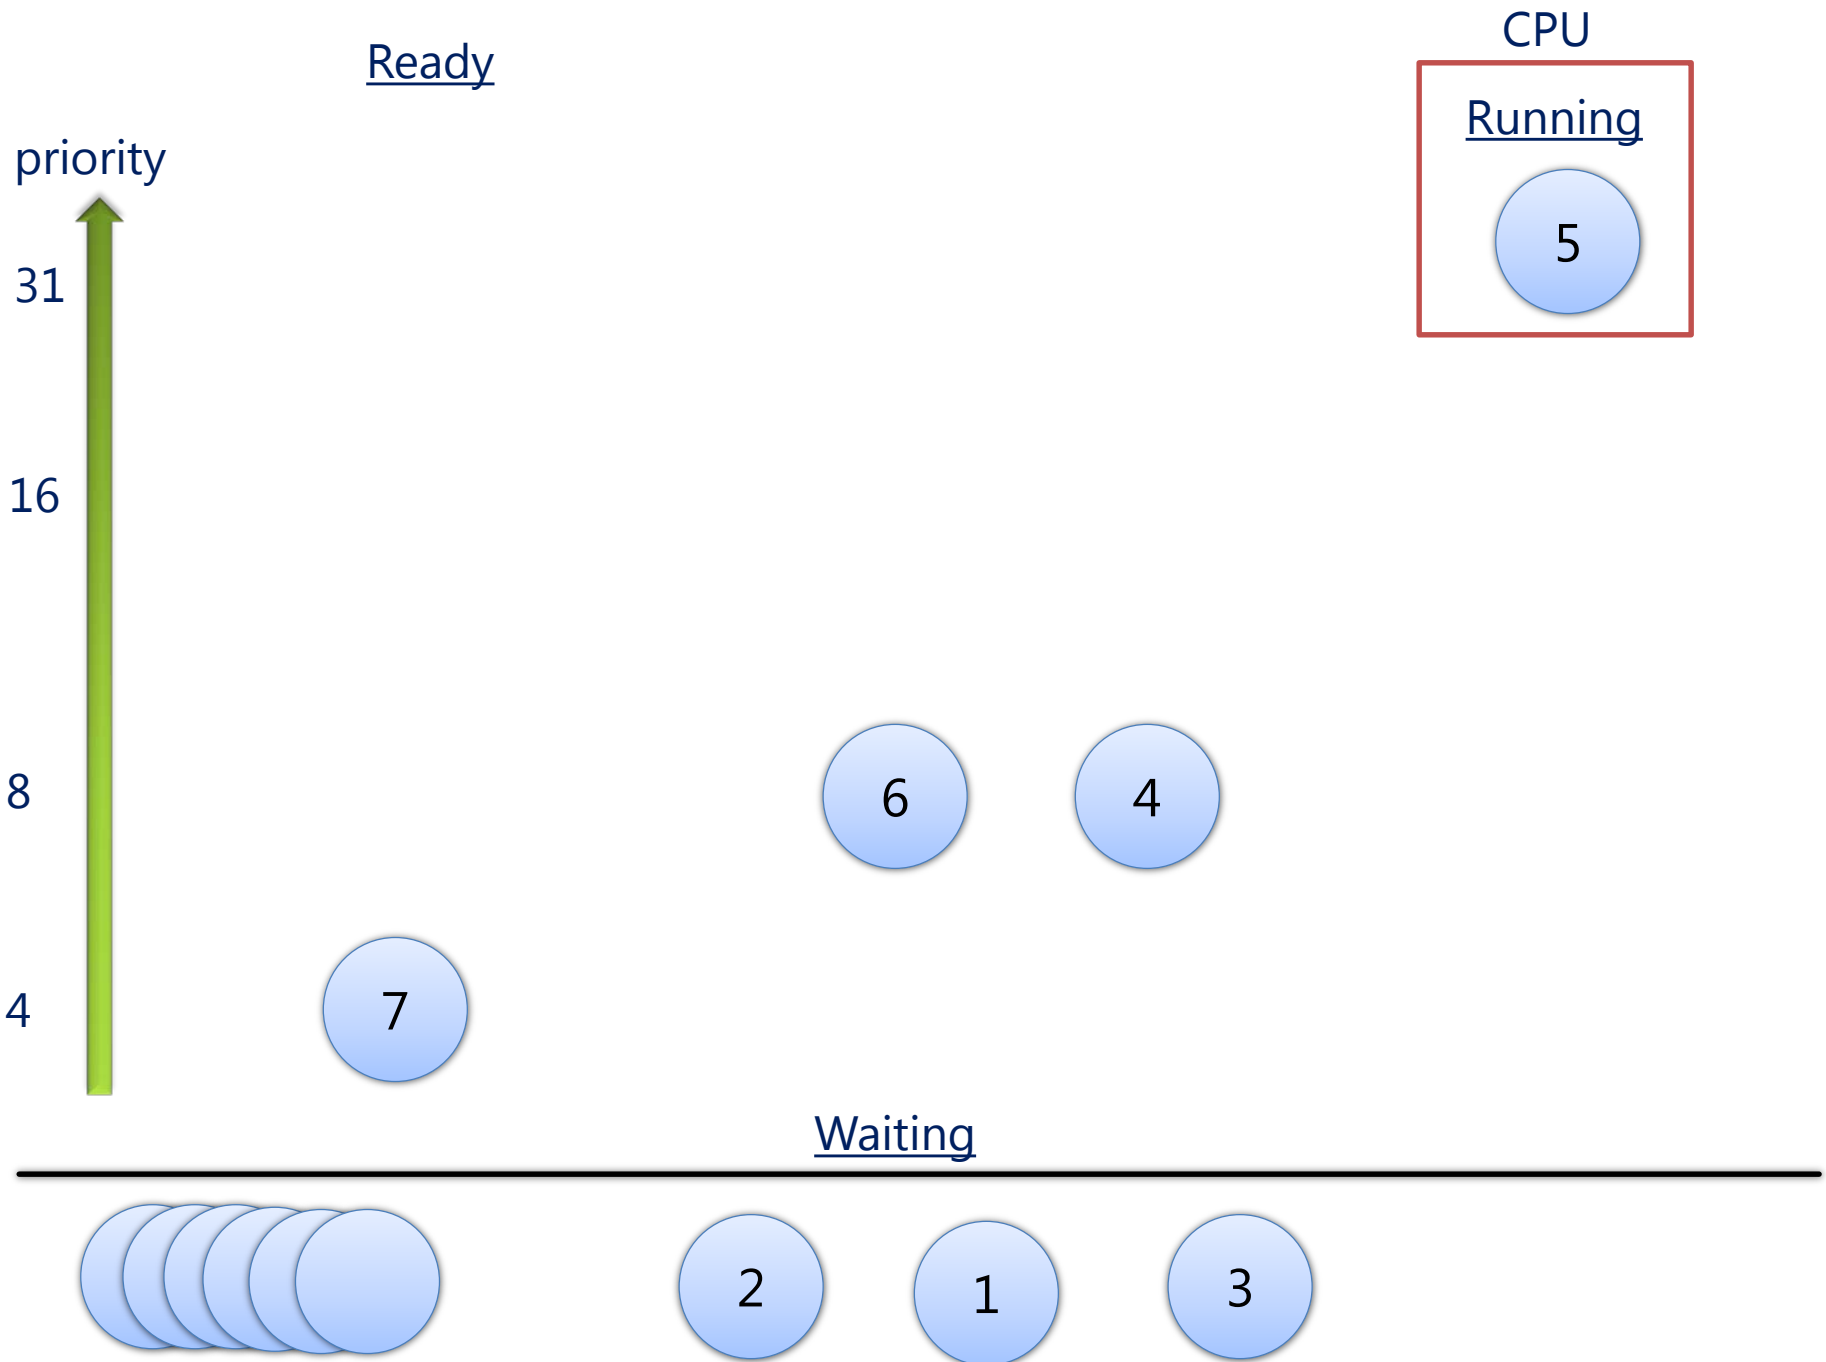
- Thread priorities are between 1 and 31 (31 being the highest)
- Priority 0 is reserved for the zero page thread
- The Windows API mandates thread priority be based on a process priority class (base priority)
- A thread's priority can be changed around the base priority
- APIs (Win32)
  - **SetPriorityClass** – changing process base priority
  - **SetThreadPriority** – change the thread priority offset from the parent's base priority
- API (kernel)
  - **KeSetPriorityThread** – change thread priority to some absolute value

# Thread Priorities (Win32 View)



**Demo**

**Thread Priorities**



# Thread Scheduling (single processor)

- **Priority based, preemptive, time-sliced**

- Highest priority thread runs first
- If time slice (quantum) elapses, and there is another thread with the same priority in the Ready state – it runs
  - Otherwise, the same thread runs again
- If thread A runs, and thread B (with a higher priority) receives something it waited upon (message, kernel object signaling, etc.), thread A is preempted and thread B becomes the Running thread

- **Voluntary switch**

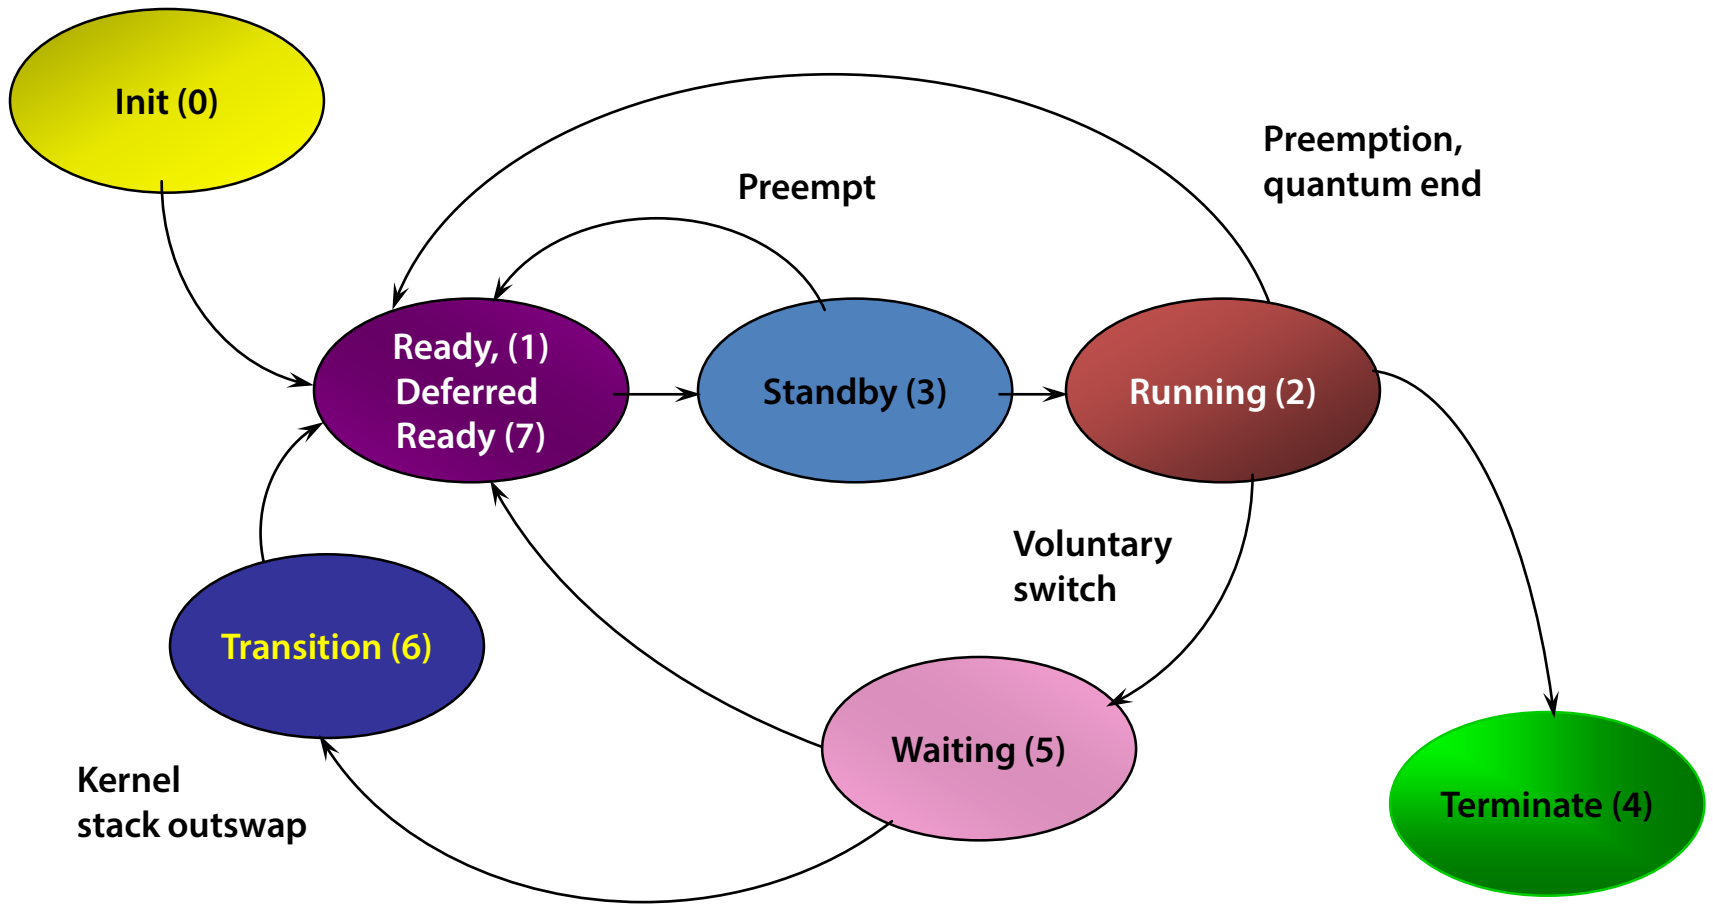
- A thread entering a wait state is dropped from the scheduler's Ready list

- **Typical time slice is 30 msec on client, 180 msec on server**

- **On an MP system with  $n$  logical processors,  $n$  concurrent threads may be running**



# Thread States



# The Scheduler

- **Scheduling routines are called when scheduling events occur**
  - Interval Timer interrupts checks for quantum end and timed wait completion
  - I/O Completion calls
  - Changes in thread priority
  - Changing state of waitable object other threads are waiting on
  - Entering a wait on one or more objects
  - Entering Sleep

**Demo**

# **Thread Scheduling**

# The Quantum

- **Scheduler clock tick is typically**
  - 10 msec (uniprocessor)
  - 15 msec (multiprocessor)
- **Can determine with clockres.exe utility from SysInternals**
- **Default client quantum is 2 clock ticks**
- **Default server quantum is 12 clock ticks**
- **Quantum can be modified by using the registry or a Job**
- **Quantum boosting**
  - On a system configured for short, variable quantum
    - The foreground process gets triple quantum
    - For any process with a priority class above Idle

# Quantum Control

- Registry key: **HKLM\SYSTEM\CCS\Control\PriorityControl**
- Value: **Win32PrioritySeparation**



- **Short vs. Long**
  - 1=long, 2=short
  - 0, 3=default (long for Server, short for Client)
- **Variable vs. Fixed**
  - 1=boost priority of foreground process, 2=don't boost
  - 0, 3=default (boost for Client, don't boost for Server)
- **Foreground quantum boost**
  - Index into a table

	Short			Long		
Variable	6	12	18	12	24	36
Fixed	18	18	18	36	36	36

**Demo**

**Thread Quantum**